

Baseline runs and final custom result

Three 89-task baseline rounds, the C0–C3 custom comparison, and the final C0-NC 89-task run.

1. Three completed baseline rounds

Each round ran the same 89 tasks with Terminus-2 and OpenHands (178 results per round). Scores below are passed tasks; the remaining tasks did not pass the official checks.

Round	Terminus-2	OpenHands
A · Original capped	18/89 (20.2%)	4/89 (4.5%)
B · API-troubled	3/89 (3.4%)	1/89 (1.1%)
C · Latest corrected	52/89 (58.4%)	44/89 (49.4%)

A used spending/billing-safety caps. B removed those caps but encountered severe model-service errors. C allowed temporary API retries and increased response/step allowances while retaining official task limits. Multiple settings changed, so score movement does not establish which change caused it.

2. Custom candidate selection (20-task development subset)

C0–C3 were separate custom variants, each scored on 20 development tasks. C0 was selected with the highest observed score; cost was not used as a tie-break.

Variant	Passed / 20	Rate	Selection
C0	15 / 20	75%	Selected
C1	14 / 20	70%	Not selected
C2	14 / 20	70%	Not selected
C3	13 / 20	65%	Not selected

C0's 15/20 is not an 89-task result. It was later revised and labelled C0-NC; see the separate full-run result on page 2.

Final custom harness: C0-NC

Candidate version 0.5.0 | One attempt per task | Full set: 89 tasks

Scope	Attempted	Passed	Zero-score failures	No verifier result
Development subset	20	15	5	0
Other 69 tasks	69	35	31	3
Full C0-NC run	89	50	36	3

50 / 89 passed · 56.2%	36 / 89 verified zero-score · 40.4%	3 / 89 setup-only; verifier not run
----------------------------------	---	---

Reading the final result

Of the 89 C0-NC attempts, 50 passed the official checks and 36 received verified zero scores. Three stopped during setup and have no verifier outcome; they are not counted as passes or zeros. The development20 and remaining69 are also shown separately above.

Context and limits

The latest corrected baseline round scored Terminus-2 52/89 and OpenHands 44/89. These are separate one-run observations; the report makes no statistical or causal winner claim. Known response-reported cost was \$1.71399138, but 799 requests had unknown cost and independent billing receipts were not verified, so total cost remains unknown. The completed C0-NC run was audited and its one off-server backup verified; full runtime restore was not tested.

Scope: C0–C3 and C0-NC results are reported separately. Later recovery attempts and matched-repeat runs are not included; diagnostic20 and confirmation60 remained deferred.